

Agent Monitor

for BES Clients

Administrator Guide

Document Revision	August 24, 2026.
Author	Mike Consuegra

Legal Notices and Trademarks

Third-Party Trademarks and Proprietary Rights

This document references products, product names, brand names, and trademarks owned by parties other than the author of this software. These references are provided for identification, interoperability, and educational purposes only. No claim of ownership, sponsorship, endorsement, or affiliation with these third parties is made or implied by their inclusion.

Product / Mark	Owner	Nature of Reference in This Document
HCL BigFix®, BigFix®, BigFix Platform™	HCLSoftware	Agent Monitor is designed to integrate with HCL BigFix via its published REST API and via standard BigFix Fixlet deployment. No source code from HCL BigFix is included in or distributed with this product.
IBM®	International Business Machines Corporation	Any references to IBM pertain to the historical evolution of the BigFix platform under prior ownership. No present-day affiliation with IBM is claimed.
PostgreSQL®	The PostgreSQL Global Development Group	PostgreSQL is an open-source relational database management system. No affiliation with or endorsement by the PostgreSQL Global Development Group is claimed.
Ubuntu®, Canonical®	Canonical Ltd.	Referenced as the operating system platform on which the Agent Monitor dashboard appliance is deployed. No affiliation with or endorsement by Canonical Ltd. is claimed.
Nginx™	F5, Inc.	Referenced as the reverse proxy component in the dashboard deployment architecture. No affiliation with or endorsement by F5, Inc. is claimed.
Python™	Python Software Foundation	Referenced as the programming language and runtime used to build the dashboard server and the Health Agent. No affiliation with or endorsement by the Python Software Foundation is claimed.
Windows®, Microsoft®	Microsoft Corporation	Referenced in the context of the Health Agent component, which runs as a Windows scheduled task on Windows endpoints. No affiliation with or endorsement by Microsoft Corporation is claimed.

General Disclaimer

All trademarks, registered trademarks, service marks, and product names mentioned in this document are the property of their respective owners. Where marks appear without the ® or ™ symbol, it shouldn't be interpreted as an indication that the mark is not protected.

Agent Monitor for BES Clients is an independent tool and **is not a product of HCLSoftware**, IBM, Microsoft, Canonical, F5, the PostgreSQL Global Development Group, the Python Software Foundation, or any other party referenced in this document.

® indicates a registered trademark. ™ indicates a trademark. All other marks are the property of their respective owners.

Table of Contents

Legal Notices and Trademarks.....	2
Third-Party Trademarks and Proprietary Rights.....	2
General Disclaimer.....	3
Table of Contents.....	4
1. Introduction.....	6
1.1 Purpose of This Guide.....	6
1.2 Intended Audience.....	6
1.3 Document Conventions.....	6
1.4 Application Overview.....	7
2. Architecture Overview.....	8
2.1 System Architecture.....	8
2.2 Component Descriptions.....	9
2.3 Network Requirements.....	10
2.4 Deployment Model.....	10
3. Getting Started: First Login.....	11
3.1 Accessing the Web Interface.....	11
3.2 Self-Signed Certificate Warning.....	11
3.3 Login Page.....	11
3.4 First-Time View.....	12
4. Core Concepts.....	13
4.1 The Health Model: Exception-Only Reporting.....	13
4.2 Health States.....	13
4.3 Watchdog, Self-Check, and Heartbeat.....	14
4.4 Criticality.....	14
4.5 Tenants.....	14
5. Dashboard.....	15

5.1 Health State Cards.....	16
5.2 Activity and SLA Cards.....	16
5.3 Charts.....	17
5.4 Tenants Table.....	18
5.5 Forwarder Liveness Strip.....	19
5.6 Refresh, Pause, and Export.....	19
6. Endpoints.....	20
6.1 The Toolbar.....	22
6.2 Reading the Table.....	22
6.3 Badges and States.....	23
6.4 Searching, Filtering, and Sorting.....	24
6.5 Pagination.....	25
6.6 Bulk Actions (Registered Endpoints).....	25
6.7 The Endpoint Detail View.....	26
7. Priority Alerts.....	28
7.1 What the Tab Lists.....	29
7.2 The Navigation Badge.....	29
7.3 Controls and Actions.....	29
7.4 How Notifications Work.....	30
7.5 Empty State.....	30
8. Administration.....	31
8.1 User Management.....	31
8.2 System Settings.....	33
8.3 Logs.....	49
9. Glossary.....	51
10. Support and Resources.....	52

1. Introduction

1.1 Purpose of This Guide

This guide explains how to use **Agent Monitor** from its web dashboard. It is a companion to the *Installation Guide*: where the Installation Guide covers standing the product up, this guide covers using it day to day, reading the health of your environment, responding to problems, and configuring the platform.

The guide is organized to follow the dashboard's navigation. After the introductory and conceptual chapters, each chapter maps to one tab in the dashboard, in the same left-to-right order you see on screen: **Dashboard**, **Endpoints**, **Priority Alerts**, and **Administration**.

1.2 Intended Audience

This guide is written for the people that run Agent Monitor: BigFix operators who watch endpoint health day to day, and the administrators who configure the platform, manage users, and set alerting policy. It assumes you are comfortable with BigFix concepts (the BES Client, operators, relays, etc.), basic networking, and Linux fundamentals for the dashboard host. It does not assume prior experience with Agent Monitor itself.

If the task requires admin-level access, the guide will say so. Most of the Dashboard, Endpoints, and Priority Alerts tabs are usable by day-to-day operators; the Administration tab is where admin privileges matter most.

1.3 Document Conventions

This guide uses the notations. The same labels appear in the callout boxes throughout the document.

Convention	Meaning
 NOTE	Additional context, clarification, or technical background.
 TIP	A helpful suggestion or best practice.
 IMPORTANT	A requirement to understand before you proceed.
 WARNING	An action that can cause data loss or a security risk.
 CRITICAL	A severe risk; take immediate action before proceeding.
Bold	On-screen labels exactly as they appear: buttons, tabs, column headers, and field names.
<code>Monospace</code>	Commands, file paths, URLs, and configuration values.

Procedures are written as numbered steps. A [SCREENSHOT: ...] marker shows where a screen capture will appear in the published guide.

1.4 Application Overview

Agent Monitor answers a single operational question that BigFix alone doesn't answer directly: *is the BES Client on each of my endpoints actually healthy and reporting?* A BigFix endpoint can appear in the console yet have a BES Client that has silently stopped checking in, stalled, or fallen out of sync. Agent Monitor places a lightweight **Health Agent** on each endpoint that watches the local BES Client and reports its health to a central **dashboard**, so that a stopped or hung client shows as a clear, actionable state.


The product has three parts:

- A **Health Agent** that runs on each BigFix-managed endpoint, monitors the local BES Client, and reports health.
- A **Dashboard** server that receives those reports, evaluates health, and presents it to operators through the web interface this guide describes.
- An optional **DMZ Forwarder** that relays reports from endpoints that cannot reach the dashboard directly (for example, endpoints served by an internet-facing relay).

Agent Monitor is deliberately quiet. Rather than a constant stream of messages, agents report by exception: when they first enroll, when an endpoint crosses into or out of an unhealthy state, and once a day as a still-alive heartbeat. This keeps the dashboard signal high and is the single most important idea to understand before reading the rest of this guide; it is covered in [Core Concepts](#).

2. Architecture Overview

2.1 System Architecture



Note: The Health Agent reports its own view of endpoint health. The dashboard independently cross-checks each endpoint against the BigFix REST API, which is how it can flag a **Disagreement**, a case where BigFix shows recent activity for an endpoint but the Health Agent has gone silent.

2.2 Component Descriptions

Component	What it does
Health Agent	Runs on each endpoint (Windows Scheduled Tasks / Linux systemd timer) with a watchdog. Monitors the local BES Client, evaluates health, and reports to the dashboard by exception. Enrolls automatically from a signed charter. Each endpoint holds its own key and encrypts its reports.
DMZ Forwarder	Optional. A relay that accepts agent reports on <code>tcp/31313</code> in an exposed network and forwards them to the dashboard. It never reads or changes report contents; it authenticates itself to the dashboard with its own key.
Dashboard application	A web application that receives and stores reports, evaluates health, cross-checks BigFix, serves the web UI, and sends alerts.
PostgreSQL	Stores endpoints, events, health history, tenants, users, and settings.
Nginx	Terminates TLS and reverse-proxies two planes: the operator web UI on <code>443</code> and the agent/forwarder ingest on <code>31313</code> .
BigFix Root Server	Not part of Agent Monitor. The dashboard queries its REST API (<code>52311</code>) read-only to cross-check endpoint activity.

2.3 Network Requirements

Port	Direction	Purpose
443 (HTTPS)	Operators → dashboard	Web user interface.
31313 (HTTPS)	Agents and forwarders → dashboard	Report ingest (the platform port). This is a separate plane from the web UI, with its own certificate.
52311 (HTTPS)	Dashboard → BigFix Root Server	Read-only REST cross-check of endpoint activity.
5001 (internal)	localhost only	The application process behind Nginx. [verify: internal app port]
5432 (internal)	localhost only	PostgreSQL.

Note: 443 and 31313 are intentionally distinct channels and use separate certificates. Operators reach the UI on 443; endpoints and forwarders never use 443, they report only on 31313.

2.4 Deployment Model

Agent Monitor runs as a single dashboard instance per environment. The application runs as one Unicorn worker (its internal scheduler holds state in memory, so it is not run multi-worker). It is multi-tenant: a single dashboard can hold many tenants, and visibility is scoped per tenant so operators see only the tenants they are entitled to. Standing up the server, ports, and certificates is covered in the *Installation Guide*.

3. Getting Started: First Login

3.1 Accessing the Web Interface

Open a browser to the dashboard's address over HTTPS:


```
https://<dashboard-address>
```

The web interface is served on port 443. Use the reporting name or IP you chose during installation.

3.2 Self-Signed Certificate Warning

If the dashboard uses the self-signed certificate generated at installation, your browser shows a one-time security warning the first time you connect. This is expected. Proceed through the warning to reach the login page. To remove the warning, replace the certificate with one your organization's browsers already trust (see the *Installation Guide*, and later the Administration chapter of this guide).

3.3 Login Page



Sign in with the administrator credentials established during installation.


► **Tip:** Change the initial administrator password immediately after first login, and create individual named accounts for each operator rather than sharing one login. User management is covered in the Administration chapter.

3.4 First-Time View

On a brand-new environment, before any endpoint has enrolled, the Dashboard tab appears with its cards at zero and its charts showing guidance rather than data. This is normal. The OS distribution panel, for example, shows:

No endpoints yet. Create an API key under **Administration** → **System Settings** → **API Keys** and start sending heartbeats.

Once endpoints begin enrolling and reporting, the cards and charts populate automatically. Enrolling endpoints is covered in the *Installation Guide*.



4. Core Concepts

These concepts underpin every tab.

4.1 The Health Model: Exception-Only Reporting

Agent Monitor is built around **exception-only reporting**. A Health Agent does not stream a constant "everything is fine" signal. Instead it reports only when something meaningful happens:

- when it first **registers** (enrolls) an endpoint,
- when the endpoint **transitions into or out of an unhealthy state**, and
- once a day as a **still-alive heartbeat**, so a healthy endpoint still proves it is present.

The practical consequence: **a quiet agent is a healthy agent**. You should not expect continuous traffic from every endpoint. What matters is that the daily heartbeat keeps arriving and that transitions are reported when they occur. This is why the dashboard emphasizes exceptions (unhealthy, silent, disagreement) rather than raw message volume.

4.2 Health States

Every endpoint is in one of a small number of states. These same terms appear on the Dashboard cards, in the Endpoints table, and in Priority Alerts.

State	Meaning
Healthy	The endpoint is reporting normally and its BES Client is behaving.
Unhealthy	The Health Agent has determined the BES Client is not reporting normally (for example, stopped or stalled).
Silent	No heartbeat has arrived from the agent within the last 24 hours.
Disagreement	The agent shows Silent, but the BigFix Console shows recent activity for the endpoint. This may mean there is a problem with the Health Agent or its path to the dashboard rather than the endpoint being down.

 **Note:** Silent and Unhealthy are different problems. **Unhealthy** means the agent is reporting and telling you the BES Client is unwell. **Silent** means the agent itself has stopped reporting, so you have lost visibility and should find out why.

4.3 Watchdog, Self-Check, and Heartbeat

An agent watches the local BES Client continuously and reports only when something meaningful changes. Three mechanisms work together:

- **The watchdog** is the primary, always-on monitor. On Linux it is event-driven: it follows the BES Client service and notices a stop or start within about a second. On Windows it polls the service on a short cadence (60 seconds by default). The moment it sees a change, it triggers an immediate report, so an outage surfaces on the dashboard in near real time.
- **The self-check** is a periodic backstop. On a fixed interval (about ten minutes by default) the agent re-evaluates the local BES Client together with host conditions such as disk, memory, clock, and its own cryptographic health, and confirms it is still alive. It catches anything the watchdog might miss and sends the daily still-alive beat. Most self-checks change nothing and therefore send nothing, this is the exception-only model at work.
- **The heartbeat** is any message that actually reaches the dashboard: on registration, on a health transition (caught by the watchdog or a self-check), and once a day as the still-alive beat.

So an agent observes far more often than it heartbeats. The dashboard's silent state is about missing heartbeats, not missing checks: a healthy agent with nothing to report still checks in once a day.

4.4 Criticality

Endpoints can be marked **critical** (priority). Criticality is what separates "worth knowing about" from "wake someone up." The Priority Alerts tab and the Dashboard's **CRITICAL UNHEALTHY** card both key off it: an ordinary endpoint going unhealthy is visible on the dashboard, but a *critical* endpoint going unhealthy is what raises an alert. Criticality is covered where it is set in the Administration chapter, and where it drives alerts in the Priority Alerts chapter.

4.5 Tenants


A single dashboard is multi-tenant: endpoints are grouped into tenants, and most views can be filtered or broken down per tenant. How an endpoint's tenant is decided:

You do not assign tenants by hand. A tenant corresponds to a BigFix environment, identified by its root server. When an endpoint enrolls, its Health Agent reads the local BES Client's actionsite masthead (the actionsite.afxm file) and reports the gather URL of its BigFix root server. The dashboard groups endpoints by that root server: one tenant per BigFix environment, created automatically the first time an endpoint from a new environment enrolls. Endpoints managed by the same root server share a tenant, and an endpoint's tenant is fixed at enrollment, so it cannot move itself to another tenant. Visibility is scoped: operators see only the tenants they are entitled to, and several settings (the silence threshold and the tenant's display name) are per-tenant. A tenant's name is derived from its root-server URL and can be given a private, per-operator alias from the Dashboard tab. Throughout this guide, "your environment" means the tenants you can see.

5. Dashboard

The **Dashboard** tab is the default landing view and the operator's at-a-glance answer to "is my environment healthy right now?" It presents, top to bottom: four verdict-count cards, four activity and SLA cards, an optional forwarder-liveness strip, a 7-day health trend, an OS-distribution donut, a Health Agent version breakdown, and a per-tenant health table.

Most items on this tab are **drill-downs**. Clicking the item takes you to the Endpoints tab already filtered to the matching population, so you move from *how many* to *which ones* in a single click. The view refreshes on a selectable interval and can be exported to PDF as a point-in-time snapshot.



5.1 Health State Cards

The top row summarizes the current health verdict for every endpoint you can see. Each card is clickable and lands you on the matching filtered Endpoints view.

Card	Sub-text	What it counts	Clicking it opens
UNHEALTHY	"BES Client not reporting normally"	Endpoints whose latest verdict is unhealthy	Endpoints → Unhealthy
SILENT	"No heartbeat in at least Nh"	Endpoints with no heartbeat inside the silence threshold (default 24h)	Endpoints → Registered Endpoints , status = silent
DISAGREEMENT	"BigFix sees activity, Health Agent silent"	Endpoints where BigFix shows activity but the agent is silent	Endpoints → Registered Endpoints , status = disagreement
HEALTHY	"Reporting normally"	Endpoints reporting normally	Endpoints → Registered Endpoints , status = healthy

5.2 Activity and SLA Cards

The second row covers recent activity and a recovery metric.

Card	Sub-text	What it shows	Clicking it opens
CRITICAL UNHEALTHY	"priority endpoints currently unhealthy"	Count of <i>critical</i> endpoints currently unhealthy (highlighted when non-zero)	Endpoints → Registered Endpoints , status = unhealthy + priority = critical
MAINTENANCE WINDOWS	"n endpoints in maintenance"	Maintenance Windows created (active and expired)	System Settings → Administration → Maintenance Windows
TOTAL EVENTS	"in the last 24 hours"	All events (heartbeats, registrations, status changes) in the last 24 hours	Endpoints → Recent Activity , last 24h

MEAN TIME TO RECOVER"over N episode(s) in the last 7 days"


The 7-day average time for an unhealthy endpoint to return to healthy; shows "—" when there were no episodes

Not clickable (display only)



► **Tip: CRITICAL UNHEALTHY** is the card to watch first. It is the same population that drives Priority Alerts, so a non-zero value here is your cue to check the Priority Alerts tab.

5.3 Charts


Endpoint Health · Last 7 Days. A horizontal timeline of 168 hourly buckets, rendered once per tenant, with each hour colored by the share of unhealthy endpoints in that hour. The legend bands are: 0% / ≤10% / ≤30% / ≤60% / >60% unhealthy, plus a gray **NO DATA** band for hours before any endpoint existed. Day boundaries are marked with tick labels and dividers at local midnight. Hovering an hour shows the unhealthy / unknown / healthy / total breakdown for that bucket. This strip loads independently, so the rest of the page appears first.



OS Distribution. A donut of endpoints by operating system family, with the total in the center and a clickable legend (swatch, label, count, percentage). Hovering a slice fans out a per-tenant breakdown for that OS. A summary line beneath reads, for example, " N tenants · N critical (N unhealthy)." Clicking a slice or legend row opens Endpoints filtered by that OS family; clicking a tenant in the hover breakdown filters by OS family *and* tenant.



Health Agent Distribution. A horizontal bar table of endpoints by agent version, with columns **VERSION / COUNT / SHARE / %**. The latest version is highlighted; older versions and an "(unknown)" row are styled distinctly. A footer summarizes "*N* active endpoints · latest: *version*." Clicking a version row opens Endpoints filtered to that version.



► **Tip:** Use **Health Agent Distribution** to spot endpoints lagging on an old agent version, then click straight through to that filtered list to plan an upgrade.

5.4 Tenants Table


The **TENANTS** table breaks endpoints and health down per tenant, with columns **TENANT / ENDPOINTS / HEALTH SPLIT**. The health split is a stacked bar showing the healthy / unknown / unhealthy proportions, with per-segment hover tooltips. Clicking a row opens Endpoints filtered to that tenant.

Each tenant name carries a pencil **rename** button that sets a private alias ("rename for me only"): ✓ **Save**, ✕ **Cancel**, and, once aliased, ↺ **Reset** to revert to the global name. Aliases affect only your own view, not other operators'.

TENANT	ENDPOINTS	HEALTH SPLIT
PROD Environment	12	
QA / LAB Environment	1	

5.5 Forwarder Liveness Strip

When one or more DMZ Forwarders exist, a strip reads **FORWARDERS — *X* of *Y* online**, with any problems called out ("*n* silent", "*n* disabled", "*n* awaiting first heartbeat"). It is green when all forwarders are online, orange if any are disabled, and red if any are silent. **If no forwarders exist, the strip is hidden entirely** rather than shown empty.



5.6 Refresh, Pause, and Export

The tab's header carries its live-view controls:

- A **last-updated stamp** shows "updated HH:MM:SS" after each successful load, or "fetch failed" in red on error.
- A **Refresh interval** dropdown selects the auto-refresh cadence: **1 / 5 (default) / 15 / 30 / 60 minutes**. Changing it restarts the timer.
- A **Pause / Resume** button stops and restarts auto-refresh; resuming reloads immediately.
- **EXPORT PDF** renders the current cards and charts into a titled "Environment Overview" snapshot sheet for sharing or record-keeping.

Note: Auto-refresh is torn down when you leave the Dashboard tab and re-established when you return, so the dashboard does not keep polling in the background while you work elsewhere. On very large environments the 7-day trend can take a few seconds to fill in; the cards and other panels appear first.

Empty and error states. Before any endpoints exist, the panels show guidance (for example, the API-key prompt described in [First-Time View](#)); with no tenants, the Tenants table reads "No tenants configured yet." While loading, panels read "Loading..." and the stamp shows "—"; on a failed fetch, panels show their own "Failed to load..." message and the stamp reads "fetch failed."

6. Endpoints

The **Endpoints** tab is where you view and manage individual agents. It lists every BES Client that runs the Health Agent and reports to the dashboard, and for each one shows live health, enrollment and authentication state, encryption posture, forwarder path, and lifecycle. It is also where per-endpoint actions live: approving or revoking an agent, rotating its key, steering how it reports, marking it critical, and hiding or purging it.

The tab has three sub-tabs, selected across the top:

Sub-tab	What it shows
UNHEALTHY	Endpoints currently in a non-healthy state. This is the default view, your working queue.
REGISTERED ENDPOINTS	The full inventory of enrolled and known endpoints. This is the only sub-tab with row selection and bulk actions.
RECENT ACTIVITY	A chronological feed of events (registration, unhealthy, recovered, heartbeat, uninstalled).

UNHEALTHY

The screenshot shows the 'UNHEALTHY' sub-tab selected in the 'ENDPOINTS' section. The interface includes a top navigation bar with 'AGENT MONITOR FOR BES CLIENTS V2.7.2', user 'Welcome, Administrator', and links for 'MY ACCOUNT' and 'LOGOUT'. A 'SERVER STATUS' indicator is green. The main navigation bar shows 'DASHBOARD', 'ENDPOINTS' (selected), 'PRIORITY ALERTS' (1), and 'ADMINISTRATION'. The sub-tab bar has 'UNHEALTHY' (selected), 'REGISTERED ENDPOINTS', and 'RECENT ACTIVITY'. A table displays one endpoint with the following data:

HOSTNAME	OS	COMPUTER ID	TENANT	STATUS	AGENT PRIORITY	AGENT VER	FORWARDER	AUTH	LAST HEARTBEAT
WIN10LAB1	WINDOWS	4745253	PROD Environment	UNHEALTHY	CRITICAL	2.7.2.1	win22labrelay ASSIGNED	ENROLLED	8/13/2026, 10:54:03 PM 11m ago

Below the table, it shows '50 per page' and 'Showing 1-1 of 1'. On the left, there are filter sections for IDENTITY (TENANT: All), OPERATING SYSTEM (All), HEALTH (AGENT PRIORITY: All, LAST HEARTBEAT: Any), SECURITY (AUTH: All), and LIFECYCLE (AGENT VERSION: All, FORWARDER: All). A 'CLEAR FILTERS' button is at the bottom left. On the right, there are 'COLUMNS', 'EXPORT CSV', and 'EXPORT PDF' options. A status bar at the top right of the table area shows 'updated 10:55:24 PM', 'REFRESH 5 minutes', and a 'Pause' button.

REGISTERED ENDPOINTS

AGENT MONITOR
FOR BES CLIENTS V2.7.2

Welcome, Administrator

MY ACCOUNT

LOGOUT

DASHBOARD

ENDPOINTS

PRIORITY ALERTS 1

ADMINISTRATION

UNHEALTHY

REGISTERED ENDPOINTS

RECENT ACTIVITY

updated 11:08:59 PM REFRESH 5 minutes Pause

FILTERS

SEARCH

Hostname or computer

IDENTITY

TENANT

All

OPERATING SYSTEM

All

BES ROLE

All

HEALTH

STATUS

All

AGENT PRIORITY

All

LAST HEARTBEAT

Any

COLUMNS

EXPORT CSV

EXPORT PDF

13 visible 1 hidden (filter:)

HOSTNAME	OS	COMPUTER ID	STATUS	AGENT PRIORITY	AGENT VER	REGISTERED DATE	LAST HEARTBEAT	FORWARDER
openclaw	LINUX	12764263	HEALTHY	—	2.7.2.1	8/11/2026, 2:06:48 AM 2d ago	8/13/2026, 4:18:43 PM 6h ago	—
WIN10LAB1	WINDOWS	4745253	UNHEALTHY	CRITICAL	2.7.2.1	7/13/2026, 1:23:41 PM 31d ago	8/13/2026, 10:54:03 PM 14m ago	win22labrelay ASSIGNED
ubuntu2404lab2	LINUX	5558692	HEALTHY	—	2.7.2.1	7/12/2026, 11:51:06 PM 31d ago	8/13/2026, 2:07:22 PM 9h ago	—
John_G_Win11v2	WINDOWS	1089154099	HEALTHY	CRITICAL	2.5.0.5	7/9/2026, 2:47:46 PM 35d ago	8/13/2026, 7:49:46 AM 15h ago	win22labrelay
WIN22LABSERVER	WINDOWS	1081951465	HEALTHY	—	2.7.2.2	7/9/2026, 12:49:37 AM 35d ago	8/13/2026, 4:01:32 PM 7h ago	—
WIN22LABRELAY	WINDOWS	14955610	HEALTHY	RELAY	2.7.2.1	7/7/2026, 9:57:22 PM 37d ago	8/13/2026, 12:35:39 PM 10h ago	—
actionarchive	LINUX	1088035851	HEALTHY	—	2.7.2.1	6/15/2026, 10:11:19 PM 59d ago	8/13/2026, 11:14:38 AM 11h ago	—
BFROOTSVR	WINDOWS	1459079	HEALTHY	ROOT	2.7.2.1	5/8/2026, 3:39:04 PM	8/13/2026, 1:21:43 PM	—

RECENT ACTIVITY

AGENT MONITOR
FOR BES CLIENTS V2.7.2

Welcome, Administrator

MY ACCOUNT

LOGOUT

DASHBOARD

ENDPOINTS

PRIORITY ALERTS 1

ADMINISTRATION

UNHEALTHY

REGISTERED ENDPOINTS

RECENT ACTIVITY

updated 11:10:54 PM REFRESH 5 minutes Pause

FILTERS

SEARCH

Hostname or computer

IDENTITY

TENANT

All

OPERATING SYSTEM

All

EVENT

EVENT TYPES

All

STATUS (AT EVENT)

All

RECEIVED

All time

STATE

COLUMNS

EXPORT PDF

Showing 30 most recent events

HOSTNAME	OS	COMPUTER ID	STATUS (AT EVENT)	AGENT PRIORITY	AGENT VER	EVENT	DETAILS	RECEIVED	FORWARDER
WIN10LAB1	WINDOWS	4745253	UNHEALTHY	CRITICAL	2.7.2.1	UNHEALTHY	BESClient service not running	8/13/2026, 10:54:03 PM 16m ago	win22labrelay
openclaw	LINUX	12764263	HEALTHY	—	2.7.2.1	HEARTBEAT	BESClient running; last BES log entry at 8/13/2026 4:08:28 PM	8/13/2026, 4:18:43 PM 6h ago	—
WIN22LABSERVER	WINDOWS	1081951465	HEALTHY	—	2.7.2.2	HEARTBEAT	BESClient running; last BES log entry at 8/13/2026 12:58:18 PM	8/13/2026, 4:01:32 PM 7h ago	—
ubuntu2404lab2	LINUX	5558692	HEALTHY	—	2.7.2.1	HEARTBEAT	BESClient running; last BES log entry at 8/13/2026 1:45:42 PM	8/13/2026, 2:07:22 PM 9h ago	—
BFROOTSVR	WINDOWS	1459079	HEALTHY	CRITICAL	2.7.2.1	HEARTBEAT	BESClient running; last BES log entry at 8/13/2026 1:16:59 PM	8/13/2026, 1:21:43 PM 9h ago	—
WIN22LABRELAY	WINDOWS	14955610	HEALTHY	CRITICAL	2.7.2.1	HEARTBEAT	BESClient running; last BES log entry at 8/13/2026 9:32:48 AM	8/13/2026, 12:35:39 PM 10h ago	—
WIN10LAB1	WINDOWS	4745253	HEALTHY	CRITICAL	2.7.2.1	HEARTBEAT	BESClient running; last BES log entry at 8/13/2026 12:21:39 PM	8/13/2026, 12:35:38 PM 10h ago	win22labrelay

6.1 The Toolbar

All three sub-tabs share an auto-refresh bar: a **last-updated** stamp, a **Refresh** interval dropdown (**1 / 5 (default) / 15 / 30 / 60 minutes**), and a **Pause** button. The timer applies to whichever sub-tab is active, and switching sub-tabs does not reset it. Each sub-tab also has a **COLUMNS** button (show, hide, and reorder columns; preferences are saved per browser) and CSV / **EXPORT PDF** actions.

6.2 Reading the Table

Columns vary slightly by sub-tab; the common ones are below. Columns are resizable (drag the edge, or double-click to auto-fit) and most are sortable (click the header to toggle ascending or descending).

Column	Meaning
HOSTNAME	The endpoint's hostname.
OS	Operating-system family badge (Windows / Linux).
COMPUTER ID	The BigFix Computer ID, or a dash if not known.
TENANT	The owning tenant; the tooltip shows its gather URL.
STATUS	The health state badge, plus a staleness indicator and any inline condition pills (see Badges and States).
AGENT PRIORITY	Criticality and BES-role indicator (see below).
AGENT VER	The Health Agent version from the last heartbeat.
FORWARDER	The forwarder that relayed the last report; a dash means the endpoint reports directly. An ASSIGNED badge means an administrator manually assigned the forwarder.
AUTH	Enrollment / authentication state (see below).
LAST HEARTBEAT	When the last heartbeat arrived (absolute date and relative time).

The **REGISTERED ENDPOINTS** sub-tab adds **REPORTING MODE** (encryption posture), **ANCHOR** (key-delivery verification), and **REGISTERED DATE**, and a leading checkbox column for bulk selection. The **RECENT ACTIVITY** sub-tab replaces STATUS with **STATUS (AT EVENT)** (the verdict at the moment of the event, not the current one) and adds **EVENT** (the event type) and **DETAILS** (the verdict reasons).

6.3 Badges and States

Agent Monitor packs a lot of an endpoint's condition into compact badges. This section is the key to reading them.

Health verdict (in STATUS): **HEALTHY**, **UNHEALTHY**, **SILENT**, **DISAGREEMENT**, and occasionally **CRYPTO_BROKEN** or **UNKNOWN**. These are the states defined in [Core Concepts](#). Next to a **HEALTHY** badge, a small **staleness pie** depletes from green to amber to grey as the endpoint approaches its silence threshold; its tooltip reads "Last activity: X ago. Will flag SILENT in N." 

Agent priority (AGENT PRIORITY column):

Badge	Meaning
CRITICAL	Marked critical by an administrator. Critical endpoints are the ones that raise Priority Alerts.
ROOT / RELAY / WEBUI / WRPT	The endpoint is a BigFix infrastructure component (root server, relay, WebUI, or Web Reports). A green "+" variant (e.g. ROOT+) means it fills more than one role. BES components are always treated as critical and cannot be downgraded.
—	Standard priority.

Authentication (AUTH column):

Badge	Meaning
ENROLLED	Approved; the agent holds its own signing key and signs every request.
PENDING	Enrolled but awaiting administrator approval.
REVOKED	An administrator cut the agent off; a reinstall is required to return. The row is also shown in red.

Reporting mode (encryption posture, Registered sub-tab): **ENCRYPTED** (compliant), **PLAINTEXT** (not encrypting, where policy allows it), **AT RISK** (policy requires encryption but the agent is not encrypting), **PENDING** (policy changed or no report yet), or **OFF** (encryption not required by policy).

Anchor (Strongbox key-delivery verification, Registered sub-tab): **ENFORCED**, **READY**, **MISMATCH**, **NOT SET**, or a dash for agents that do not report it.

Inline condition pills can appear beside the STATUS badge:

Pill	Meaning
UNINSTALLED	The dashboard received an uninstall/goodbye event for this endpoint.
CRYPTO BROKEN	The agent reports a damaged cryptography install. Recovery is a reinstall with an antivirus exclusion; the pill clears automatically on the next signed heartbeat.
TAMPERED	The reported binary's SHA-256 does not match the build attestation on file.
CHAIN INVALID	(Windows) The agent's Authenticode signature does not chain to the Agent Monitor root certificate on the endpoint.
UNKNOWN BUILD	No attestation is on file yet for that version and platform.

Note: A clean, correctly attested agent shows no integrity pill at all, the healthy state is silent, exactly as on the Endpoints list and in the Installation Guide.

6.4 Searching, Filtering, and Sorting

Each sub-tab has a collapsible **FILTERS** rail on the left with a **CLEAR FILTERS** button. Filters are grouped under headings (Identity, Health, Security, Lifecycle, Event). Active filters appear as chips above the table, and the current view is reflected in the page URL, so a refresh or a bookmark restores the same filtered view.

Common filters include **SEARCH** (hostname or computer ID), **TENANT**, **OPERATING SYSTEM**, **AGENT PRIORITY** (Critical / Non-critical), **LAST HEARTBEAT** (within 1 hour / 24 hours / 7 days / no heartbeat in 24h+), **AUTH**, **AGENT VERSION**, and **FORWARDER** (Direct or a specific forwarder). The **REGISTERED ENDPOINTS** rail adds **STATUS** (Healthy / Unhealthy / Silent / Disagreement / Crypto Broken), **BES ROLE**, **REPORTING MODE**, **ANCHOR**, **REGISTERED** (age), and **VISIBILITY** (Visible / Hidden / All). The **RECENT ACTIVITY** rail adds **EVENT TYPES** (a multi-select of Registration / Unhealthy / Recovered / Heartbeat / Uninstalled), **STATUS (AT EVENT)**, and **RECEIVED** (time window).

To sort, click any sortable column header; click again to reverse it.

6.5 Pagination

On the UNHEALTHY, REGISTERED and RECENT ACTIVITY sub-tabs, results are paged. A per-page selector offers **50 / 100 / 250 / 500 / 1000**, a status line reads "Showing X–Y of Z," and first / previous / next / last controls move between pages.

6.6 Bulk Actions (Registered Endpoints)

The **REGISTERED ENDPOINTS** sub-tab lets you act on many endpoints at once. Select rows with the per-row checkboxes or the header "select all on this page" box. Selection is scoped to the current page and clears whenever the list reloads (paging, filtering, sorting, or switching sub-tabs).

While at least one row is selected, a bar shows "*N* selected," a **Clear** link, and an **Actions ▼** menu. Menu items enable or disable themselves based on the selection, with a tooltip explaining any disabled item:

Group	Actions
Forwarder	<i>Assign to forwarder...</i> (manually assign a forwarder, overriding automatic selection) · <i>Clear forwarder assignment</i> (return to automatic).
Groups	<i>Add to group...</i> · <i>Remove from group...</i> (with inline + NEW GROUP creation when adding).
Authentication	<i>Revoke enrollment</i> (requires a reason; cuts the agents off) · <i>Restore enrollment</i> (un-revoke, same key). Mixed selections disable both.
Hidden agents	<i>Hide selected...</i> (reversible; a heartbeat restores it) · <i>Un-hide selected...</i> · <i>Purge selected...</i> (permanent; requires typing PURGE and lists the exact hostnames).
Critical Priority	Assign or clear Critical Priority from an agent.



Warning: *Purge* permanently removes hidden endpoints and their history. It requires a typed confirmation and cannot be undone. *Hide* is the reversible alternative, a hidden endpoint returns on its next heartbeat.

6.7 The Endpoint Detail View

Click any row to open the endpoint detail view, titled with the hostname. It is both the fullest picture of one endpoint and the place most per-endpoint administrator actions live (the image below is a partial view).

The screenshot displays the 'UBUNTU2404LAB2' endpoint detail view. The left sidebar contains sections for Metadata, Authentication, and Report Encryption. The main content area on the right includes a Status Trend chart, Episode statistics, a Diagnostic Bundle with various system metrics, Relay Reachability status, BES Client configuration and settings, Network State, and a Log Tail showing system logs.

The detail view also includes a **Report Encryption** summary (policy mode, the last report's wire form, and key id), a **Status Trend** chart with selectable windows (**1H / 4H / 1d / 7d / 30d / 90d**) and episode statistics, a **Diagnostic Bundle** section with the latest collected diagnostics (relay reachability, BES Client config and settings, network state, communication errors, and a log tail), and a **Last 20 Events** table.

Note: The Diagnostic Bundle is a deeper troubleshooting surface. When an endpoint is unhealthy, it is usually the fastest way to see *why*, the relay reachability checks and communication-error entries in particular.

At the top, conditional banners flag special conditions: **CRYPTO BROKEN** (with the reported error and recovery guidance), **HIDDEN ENDPOINT** (soft-deleted, with who and when), or **HIDDEN VIA TENANT** (the endpoint's tenant was hidden).

Metadata covers identity and health: hostname, tenant, OS and version, observed IP address (and the agent-reported IP when it differs), BES Client version, BigFix Computer ID, Health Agent version, registration and last-activity times, the status at last report (with reasons), mean recovery time, group memberships (as clickable chips), silent-since and uninstalled timestamps, and read-only notes. It also holds the **Agent Priority** critical checkbox (administrator-only; locked for BES components).

Authentication shows enrollment status (including **PERMANENT** or **ROTATION PENDING** labels), the tenant's enrollment mode, the signature algorithm (for example `ed25519` or `m1-dsa-65`), the public-key fingerprint, enrollment and approval times, last key rotation, revocation details, and the report-routing state.

The administrator actions live here as well. They are available to administrators only, and each appears according to the endpoint's current state, so you will not see them all at once. They are grouped below by when they apply.

Action	Effect
Rotate key	Approved endpoints. Issues a new same-algorithm keypair on the next heartbeat. While a rotation is pending, a Cancel rotation control replaces it.
Revoke enrollment	Approved endpoints. Cuts the agent off. Requires a typed reason, which is recorded in the audit log.
Restore	Revoked endpoints, before Make permanent. Un-revokes the endpoint, reinstating its existing key.
Make permanent	Revoked endpoints, before Make permanent. An irreversible escalation; recovery then requires a reinstall with a new identity. Afterward Restore and Make permanent both disappear, replaced by a PERMANENT label.
Route via forwarder / Force dashboard / Clear	Any enrollment state. Steer how the endpoint reports, with no Fixlet; the change takes effect on the next heartbeat acknowledgement. Clear appears only when routing is already dashboard-managed. The current path shows as a “via forwarder X” or “reporting direct” pill.
Hide endpoint / Un-hide / Purge now	Visibility controls. Hide endpoint when the endpoint is visible; Un-hide and Purge now when it is already hidden. See Bulk Actions for the same operations in bulk.

► **Tip:** Route via forwarder and Force dashboard are the per-endpoint equivalents of the bulk forwarder actions. Use them to move a single endpoint's reporting path without touching its configuration file or issuing a Fixlet; the change takes effect on the endpoint's next heartbeat.

7. Priority Alerts

The **Priority Alerts** tab is a live worklist of endpoints that need attention right now. Its scope is deliberately narrow: it lists only alerts where an endpoint is both **critical** and **unhealthy**. An alert is raised the moment a critical endpoint transitions from healthy to unhealthy, and it clears automatically the moment that endpoint reports healthy again. Cleared alerts are not shown, so the tab always reflects the present, not history.

This is a monitoring and triage surface, not a configuration surface. Who gets emailed, and for which endpoints, is defined under Administration (on endpoint groups); this tab shows only the *effect* of that configuration.

Note: This is the reason **criticality** matters (see [Core Concepts](#)). An ordinary endpoint going unhealthy is visible on the Endpoints tab but does not raise a Priority Alert. Only a *critical* endpoint going unhealthy lands here and triggers a notification.

The screenshot shows the 'AGENT MONITOR FOR BES CLIENTS V2.7.2' interface. The top navigation bar includes 'DASHBOARD', 'ENDPOINTS', 'PRIORITY ALERTS 2', and 'ADMINISTRATION'. The 'PRIORITY ALERTS' tab is active, showing a table of alerts. The table has columns: RAISED, HOSTNAME, TENANT, CURRENT STATUS, EMAIL, and ACKNOWLEDGED. Two alerts are listed, both with 'UNHEALTHY' status. The first alert is for 'WIN10LAB1' in 'PROD Environment', with a 'SENT' email status and an 'Acknowledge' button. The second alert is for 'ubuntu2404lab2' in 'PROD Environment', with a 'PENDING' email status and an 'Acknowledge' button. The interface also includes a 'SERVER STATUS' indicator, a 'REFRESH' button, and a 'SEND ALERT NOW' button.

RAISED	HOSTNAME	TENANT	CURRENT STATUS	EMAIL	ACKNOWLEDGED
8/14/2026, 1:20:17 PM	WIN10LAB1	PROD Environment	UNHEALTHY	SENT	Acknowledge
8/14/2026, 1:19:50 PM	ubuntu2404lab2	PROD Environment	UNHEALTHY	PENDING	Acknowledge

7.1 What the Tab Lists

At the top, a summary bar shows two counts, **Active** and **Acknowledged**. Below it, one row per open critical-unhealthy alert. Clicking a row (anywhere except the Acknowledge button) opens that endpoint's detail view.

Column	Meaning
RAISED	When the alert was raised.
HOSTNAME	The endpoint's hostname.
TENANT	The owning tenant.
CURRENT STATUS	The endpoint's live verdict (for an open alert, unhealthy).
EMAIL	The notification state: SENT once the email has gone out (the tooltip names the channel), or PENDING while it waits for the next dispatch.
ACKNOWLEDGED	If someone has acknowledged the alert, shows "by <i>username</i> "; otherwise shows an Acknowledge button.

An alert is in one of three states:

- **Active:** open and unacknowledged.
- **Acknowledged:** open, but an operator has taken ownership. Acknowledging does not resolve the alert; it signals that someone is on it. The row stays visible and moves from Active to Acknowledged.
- **Cleared:** the endpoint recovered to healthy. This happens automatically and the row drops off the tab. There is no manual "resolve", recovery is the only way an alert closes.

7.2 The Navigation Badge

The number on the **PRIORITY ALERTS** navigation tab is the **Active** count, open alerts that no one has acknowledged yet. It is hidden when there are none. This badge refreshes about once a minute from anywhere in the application, so it stays current even while you are on another tab, making it your ambient "something needs me" indicator.

7.3 Controls and Actions

An auto-refresh bar carries a **last-updated** stamp, a **Refresh interval** dropdown (**15s / 30s (default) / 1m / 2m / 5m / 15m**), and a **Pause / Resume** button. A controls row holds:

- A **STATUS** filter: **All (open)** (default), **Active (unacknowledged)**, or **Acknowledged (still open)**. This is the only list filter; the critical-unhealthy scope is fixed.
- **REFRESH:** reload the list now.

- **SEND ALERT NOW:** trigger an immediate notification sweep instead of waiting for the scheduled one. A toast reports the outcome ("sent *N* alert(s)", "no pending alerts", "no recipients resolved", or that dispatch is disabled or SMTP is unconfigured). This is handy right after you change SMTP or recipient settings.
- **COLUMNS:** show, hide, and reorder columns (saved per browser).
- **EXPORT PDF:** render the current filtered view to a "Priority Alerts" sheet, stamped with the active filter.
- **Acknowledge** (per row) take ownership of an alert. It records who acknowledged it and when.

7.4 How Notifications Work

The **EMAIL** column reflects a background dispatcher that runs on a schedule (about every five minutes by default) and can also be run on demand with **SEND ALERT NOW**. Three behaviors are worth understanding:

- **Recipients come from groups.** For each alert, recipients are gathered from every group the endpoint belongs to whose "notify on unhealthy" toggle is on. An endpoint in several groups notifies all of them.
- **There is no global fallback.** If none of an endpoint's groups has recipients configured for this category, the alert stays **PENDING**, visible on the tab, but no email is sent, until an administrator assigns recipients on one of its groups. This is intentional: alerts are never sent to a catch-all address.
- **One email per recipient, bundled.** If a recipient owns ten endpoints that all go critical-unhealthy in the same sweep, they receive a single email listing all ten, not ten separate messages.

► **Tip:** A row stuck on **PENDING** almost always means the endpoint's group has no alert recipients set. Fix it under **Administration → Endpoint Groups** by adding recipients to a group the endpoint belongs to, then use **SEND ALERT NOW** to dispatch without waiting.

7.5 Empty State

When nothing needs attention, the list reads: "**No critical-unhealthy alerts. Mark endpoints critical from the endpoint detail panel to enable alerting.**" That is the intended steady state, an empty Priority Alerts tab means every critical endpoint is healthy.

8. Administration

Administration is where you configure the platform itself. It is organized into three areas, shown as sub-tabs across the top: User Management, System Settings, and Logs.

Note: The entire Administration tab is available to administrators only. Auditors and viewers cannot open these pages; access is enforced on the server, not merely hidden in the interface.

8.1 User Management

User Management controls who can sign in and what they can see. User accounts are assigned to a specific role. Non-administrator accounts can be limited to the tenants administrators give them access to see.

Accounts

The Accounts page lists every dashboard account with its username, display name, email, role, scopes, status, last login, and creation date. Use the **CREATE USER** button at the top right to add an account. You set the username (User names are permanent and can't be renamed later), email, display name, role, and an initial password that the user must change at first login.

Each row offers per-account actions: **Edit**, **Scopes** (for non-admins, to choose which tenants they see), **Reset PW** (sets a new password the user must change at next login), **Unlock** (clears the failed-login counter and any active lockout), and **Delete**. You cannot deactivate, demote, delete, or bulk-change your own account. Select multiple rows to assign a role in bulk.

Note: Access is deny-by-default. A non-administrator with no tenant scopes sees nothing until you grant scopes. Deleting an account is permanent. If you will need the user account again deactivate it instead by clearing the **Account Active** checkbox in the user's Edit dialog.

USER ACCOUNTS

+ CREATE USER

☐	USERNAME	DISPLAY NAME	EMAIL	ROLE	SCOPES	STATUS	LAST LOGIN	CREATED	ACTIONS				
☐	admin	Administrator	mike.consuegra@gmail.com	ADMIN	ALL (BYPASS)	ACTIVE	8/17/2026, 9:34:24 AM	5/6/2026, 4:27:12 PM	Edit	Unlock			
☐	mconsuegra	Mike Consuegra	mxc0bbn@gmail.com	ADMIN	ALL (BYPASS)	ACTIVE	7/22/2026, 1:25:18 PM	5/12/2026, 8:03:52 PM	Edit	Reset PW	Unlock	Delete	
☐	jonsnow	King In the North	kinginthenorth@winterfell.com	ADMIN	ALL (BYPASS)	ACTIVE	—	6/16/2026, 9:37:02 PM	Edit	Reset PW	Unlock	Delete	
☐	sansa	Sansa Stark	c5z06.mike@gmail.com	VIEWER	1 TENANT	ACTIVE	8/5/2026, 1:51:23 PM	6/30/2026, 5:00:27 PM	Edit	Scopes	Reset PW	Unlock	Delete

Roles

The Roles page shows the three fixed roles and how many users hold each. Currently roles can't be created or edited; you assign a role when you create or edit an account. **Manage Members** lists the users in a role and lets you open any of them for editing.

Role	Access
Admin	Full access. Manages users, roles, settings, connections, API keys, and the audit log across all tenants.
Auditor	Read-only access plus the audit log, limited to assigned tenants. Can not change settings or users.
Viewer	Read-only operator. Dashboard and endpoint detail for assigned tenants only.

ROLE MANAGEMENT

Assign users to roles that control their access level within the application. Every user has exactly one role; change a user's role from the Accounts page or via Manage Members below. Non-admin roles are limited to their assigned tenants (set per user under Accounts → Scopes).

ADMIN

5 users

Full administrative access. Manages users, roles, settings, connections, API keys, and the audit log across all tenants.

Manage Members

AUDITOR

1 user

Read-only access plus the audit log, limited to assigned tenants. Cannot modify settings or users.

Manage Members

VIEWER

1 user

Read-only operator. Dashboard and endpoint detail for assigned tenants only.

Manage Members

8.2 System Settings

System Settings holds the platform's configuration pages. Most are inline-edit: click a value, change it, and it saves when you navigate away from the field (or press Enter); press Esc to cancel. Dropdowns and toggles save as soon as you change them. A few pages (Connections, Certificates, and Backup) use explicit forms and buttons.

Connections

Connections configures the external systems the dashboard talks to. The BigFix Root Server card is per-tenant (pick a tenant from the list): set the Server URL, Port (52311 by default), authentication (username and password, or an API token), and whether to verify the server's TLS certificate. Secrets are encrypted at rest and never sent back to the browser. Use **Test connection** to verify credentials before you **Save**.

Each tenant can also carry a BigFix Root CA certificate. **Capture from server** retrieves the server's CA for you to review and install, or you can paste it as PEM. When a Root CA is installed it takes precedence and the Verify SSL setting no longer applies. From this card you can also **Rotate all keys** for the tenant's endpoints, or **Hide tenant** (which requires a typed reason and hides the tenant and its endpoints from the interface). The Database card holds the dashboard's own PostgreSQL connection; changes there require a service restart.

► **Tip:** Always use Test connection before saving a BigFix or database change. It verifies the credentials without committing them.

Agent API Keys


API keys are the tenant-level gate at the ingestion point; every Health Agent also signs its own requests, so both must check out. The table shows each key's name, purpose, status, last-used and expiry, and who created it. Use **+ CREATE NEW** to create one, choosing a purpose:

Purpose	Use
Health Agent	Authenticates heartbeat and diagnostic ingest.
Integration	Other API access for external integrations.
Admin	Administrative API access.

❗ **Note:** A key's raw value is shown only once, at creation, and is stored only as a hash, so copy it then; it can never be retrieved again. Revoke disables a key permanently; Delete only removes the row. Forwarder keys are managed on the Forwarders page, never here.

Forwarders

Forwarders are DMZ or relay nodes that carry agent reports for endpoints with no direct path to the dashboard. Each forwarder has its own key and TLS certificate. The table groups forwarders by tenant and shows status, last heartbeat, spool depth, version, endpoint count, and certificate state. The cadence bar at the top sets how often forwarders heartbeat and how long an idle one may go before it is flagged SILENT.



Status	Meaning
ONLINE	Reporting normally within the heartbeat cadence.
SILENT	No heartbeat within the silent threshold.
DISABLED	Turned off; its key is left intact.
NEVER SEEN	Registered but not yet heard from.
REVOKED	Key permanently killed; cannot report.

Use **+ CREATE NEW** to register a forwarder (pick its tenant and name) and create its key, which is shown once but can be viewed again later with an audited reason. Per-forwarder actions include **View key** (requires a reason; every view is logged), managing its TLS certificate (dashboard-issued or bring-your-own CA), **Disable** and **Enable**, and **Revoke**. Disable stops a forwarder's relayed traffic without destroying its key; Revoke kills the key for good.

Placement (Internal or DMZ). The PLACEMENT column shows where each forwarder sits, and you set it from the dropdown in that column (it defaults to Internal). Placement tells agents when to use a forwarder. Internal forwarders are reachable only on the corporate network; DMZ forwarders are reachable from the public internet. An agent on the corporate network reports to the dashboard directly or through an internal forwarder, and dials a DMZ forwarder only when nothing internal answers, which is also how it recognizes it has moved off the network. An agent that is off the network tries to route to the dashboard first, then to a DMZ forwarder, and returns to the internal path automatically once it can reach one again.

Endpoint Groups

Endpoint groups route critical-unhealthy alerts to specific people. assign one or more email addresses to a group, add the endpoints those people own, and alerts about those endpoints email only that group's recipients. An endpoint can belong to many groups. The list shows each group's tenant, recipients, and member count; filter by tenant or search by name.

Use **+ NEW GROUP** to create one (name, an optional description, and the comma-separated alert recipients). Open a group to edit it, manage members (add by live search, or **Remove**), and see each member's status and critical flag. Removing an endpoint from a group does not delete the endpoint.

ENDPOINT GROUPS ? + NEW GROUP

FILTER BY TENANT

- All Tenants 4
- PROD Environment 2
- QA / LAB Environment 2

SEARCH
Filter group names...

NAME	TENANT	ALERT RECIPIENTS	MEMBERS	ACTIONS
Mission Critical Devices ?	PROD Environment	[REDACTED]	5	Edit Delete
All Endpoints ?	PROD Environment	NONE	12	Edit Delete
John-Test ?	QA / LAB Environment	[REDACTED]	1	Edit Delete
All Endpoints ?	QA / LAB Environment	NONE	1	Edit Delete

Note: A group with no recipients cannot deliver alerts. Alerts on its endpoints stay PENDING until you set at least one recipient.

Health Agent Configuration

This page holds the agent settings the dashboard pushes to Health Agents; each agent adopts new values on its next heartbeat, with no reinstall or Fixlet. The settings are grouped:

- **Agent cadence and diagnostics:** the self-check backstop interval, the Windows BES Client poll interval, the log-activity verdict window, whether to attach a diagnostic bundle when an endpoint is unhealthy, and the maximum log-tail size.
- **Forwarder selection:** how often an agent re-checks its forwarder choice, and the margins a candidate must beat before the agent switches to it.
- **Silent and disagreement thresholds:** how long without a heartbeat marks an endpoint SILENT, and the cross-check window for a BigFix-versus-agent disagreement.
- **Agent enrollment:** the permitted signature clock skew, and the post-quantum algorithm rotation and downgrade switches.

Note: Data retention and the BigFix cross-check cadence are configured on the Maintenance page, not here.

HEALTH AGENT CONFIGURATION

AGENT CADENCE & DIAGNOSTICS

Self-Check Backstop Interval

Self-check interval that runs as a backup to the watchdog daemon.

3600 seconds

Windows BES Client Poll Interval

How often the Windows watchdog daemon queries the BES Client service state. Lower values detect transitions faster; higher values reduce CPU and process-spawn footprint.

1 minute (default, fast) ▼

Log-Activity Verdict Window

How long the BES Client log may be silent before an endpoint is flagged unhealthy; 4x this window escalates to a hard failure. Applies to both platforms and is pushed to agents on their next heartbeat (agents fall back to their local config.ini until the first ack).

7200 seconds

Attach Diagnostic Bundle

When a Health Agent reports "unhealthy", include a diagnostic bundle (BES Client logs, system state) with the event.

Enabled ▼

Max Log Tail Size

Upper bound on each log attachment included in a diagnostic bundle. Hard maximum 2048 KB (2 MB); the dashboard rejects larger values.

512 KB

FORWARDER SELECTION

Forwarder Re-Check Interval

How often each Health Agent re-measures TCP-connect proximity across the forwarders it has been given and may switch to a closer one. Selection is sticky between re-checks. Pushed to agents on their next heartbeat.

86400 seconds

Switch Margin (percent)

Anti-flap: at a re-check, a different forwarder must be faster than the current one by at least this percentage (or the millisecond margin, whichever qualifies) before the agent switches. Prevents bouncing between two near-equal forwarders.

25 %

Switch Margin (milliseconds)

Anti-flap floor: a different forwarder must beat the current one by at least this many milliseconds (or the percentage margin) before the agent switches.

30 ms

SILENT / DISAGREEMENT THRESHOLDS

Silent Threshold

After this many hours without any signal from a Health Agent, the endpoint is flagged Silent.

24 hours

Disagreement Cross-Check Threshold

Hours an endpoint must remain Silent before it becomes a candidate for the nightly BigFix REST cross-check.

72 hours

AGENT ENROLLMENT

Signature Clock Skew

Max acceptable clock skew between a Health Agent and the dashboard for signed-request timestamps. Same window bounds the server-side replay-protection cache.

300 seconds

Report Encryption

Report Encryption (Strongbox) controls end-to-end encryption of Health Agent report payloads. Set the global Policy, and optionally override it per tenant:

Policy	Behavior
Off	No encryption. This is the default.
Optional	Encrypt where the agent is capable, but still accept plaintext. Used during a rolling migration.
Required	Reject plaintext; every endpoint must seal its reports.

REPORT ENCRYPTION ?

STRONGBOX ?

Policy
Required: reject plaintext (all endpoints m

Per-tenant overrides ?

PROD Environment

Required: reject plaintext

QA / LAB Environment

Optional: encrypt if capable, allow plainte

Key-delivery anchor ?

c30e4cd611a81fd810da43ef46856c2471a1a0c824fa904a674c51e26fde1246
COPY
ML-DSA-65

The dashboard's key-delivery anchor. Your charter already carries and records it at enrollment, so agents verify key deliveries automatically. Shown here for verification (for example, to confirm a charter carries this anchor, or to spot a change after a key rotation).

Enforcement readiness ?

All 13 agents are Ready (verifying key deliveries). Safe to enforce.

13

Ready

13

Enforced

0

Not set

0

Mismatch

0

Not reported

Current strongbox keys ?

TENANT	ACTIVE_KEY_ID	CREATED	STATUS	ACTION
	2	5/23/2026, 1:25:31 AM	ACTIVE	Rotate
	1	7/9/2026, 8:14:40 AM	ACTIVE	Rotate

EMERGENCY: INVALIDATE ALL STRONGBOX KEYS

Deletes every tenant's strongbox keys. The dashboard regenerates a fresh keypair the next time any Health Agent requests one. Use this after a server-migration if the inherited keys are unreadable ("strongbox_stale_key" errors in ingest logs that don't clear after Health Agents re-fetch), or when responding to a suspected dashboard-host compromise. All Health Agents will hit one round of 409 and re-enroll cleanly on the next post.

Invalidate all strongbox keys...

The page also publishes the key-delivery anchor (a fingerprint you deploy to agents out-of-band so they can verify key deliveries), shows enforcement readiness (how many agents are Ready, Enforced, Not set, Mismatch, or Not reported), and lists the current per-tenant keys, each with a **Rotate** action.



Warning: Do not set the policy to Required until the readiness panel shows every agent Ready. Enforcing early makes un-verified agents fail to report and appear SILENT.

✗ Critical: Invalidate all strongbox keys is a destructive, last-resort action, used after a server migration or a suspected compromise. It requires you to type INVALIDATE to confirm, and reports still spooled under the old keys become unreadable.

Security

Security settings govern how accounts authenticate. Every value is inline-edit and takes effect immediately. They are grouped into Password policy (length and complexity rules, expiration, history, and whether the admin account is exempt), Account lockout (a temporary lockout after repeated failed logins), IP rate limiting (with a whitelist), Session management, and the Login banner (an optional authorized-use message shown above the login form).

SECURITY SETTINGS

Click any value to edit. Changes save when you click away (or press Enter); press Esc to cancel.

PASSWORD POLICY

Rules enforced when users set or change their dashboard password.

Enable Password Expiration When enabled, users must change their password after a fixed number of days.	Disabled
Exempt Admin Account When yes, the primary admin account is exempt from password expiration. Enable Password Expiration to edit this.	Yes (admin never expires)
Password History Count Prevent reuse of the last N passwords. Set to 0 to disable history checking.	5 passwords
Maximum Password Age Force a password change after this many days. Only applies when expiration is enabled. Enable Password Expiration to edit this.	120 days
Minimum Length Minimum number of characters required in a password.	4 chars
Require Lowercase Password must contain at least one lowercase letter (a-z).	Disabled
Require Number Password must contain at least one digit (0-9).	Disabled
Require Special Character Password must contain at least one non-alphanumeric character.	Disabled
Require Uppercase Password must contain at least one uppercase letter (A-Z).	Disabled

ACCOUNT LOCKOUT

Temporary lockout after repeated failed logins on a single account.

Lockout Duration How long an account stays locked after exceeding the failed-attempt limit.	15 minutes
Max Failed Attempts Number of consecutive failed logins that triggers a temporary account lockout.	5 attempts

Certificates

The Certificates page shows the dashboard's current TLS certificate: its subject, issuer, validity dates (with a warning as expiry nears), fingerprint, and serial, and whether it is self-signed or CA-signed. Use **Validate & install** to upload a CA-signed certificate and its private key; the dashboard checks the pair, backs up the existing certificate, reloads the web server, and rolls back automatically if the new certificate fails to load.

Note: This certificate secures both the web interface and the report-ingest endpoint. Your browser reconnects on the new certificate at its next request.

TLS CERTIFICATES ?

STATUS	CA-SIGNED
SUBJECT	C=US, ST=State, L=City, O=Lab, OU=Infra, CN=agentmonitor.lab.com
ISSUER	C=US, O=Agent Monitor, OU=Lab PKI, CN=Agent Monitor Lab Root CA
VALID FROM	Jul 31 15:12:12 2026 GMT
VALID UNTIL	Jul 28 15:12:12 2036 GMT expires in 3625d
FINGERPRINT	35:B1:1C:B4:99:63:F6:0C:E0:08:5D:C9:D4:16:58:B4:59:D4:C6:C9
SERIAL	0B74CC7E5F344F0952669445EB7718BCC7962CCE

REPLACE CERTIFICATE

CERTIFICATE (.CRT / .PEM)

Choose FileNo file chosen

PRIVATE KEY (.KEY / .PEM)

Choose FileNo file chosen

Validate & Install

Note: A backup of the existing cert is kept in `/opt/bigfix-health-archive/backups/certs/` before the swap. After a successful install, nginx is reloaded automatically; your browser session may need to reconnect on the next request.

OTHER CERTIFICATES

This page manages the dashboard's web TLS certificate only. The agent code-signing certificates (your signing CA and signing certificate) are managed alongside the agent binaries they sign, on the [Agent Builds](#) page.

Agent Builds

Agent Builds is a read-only registry of vendor-attested, known-good Health Agent build fingerprints. Entries arrive automatically from an agent's signed attestation, or you can **Import** an attestation file; fingerprints can never be typed by hand, only vendor-signed statements are accepted. The table lists each build's version, platform, SHA-256, thumbprint, source, and when it was recorded.

This page also generates a deployment charter. Enter your dashboard ingest URL, optionally name an enrollment key, keep Enforce anchor verification on, and use **Generate & download charter** to produce a signed charter file that bundles the dashboard address, a fresh enrollment key, and the anchor for first-time agent deployment.

AGENT BUILDS ?

DEPLOYMENT: CHARTER ?

Generates a fresh enrollment key and bundles it with this dashboard's address and anchor, signed by the anchor. Hand it out next to the generic installer, the endpoint verifies it and lands here automatically.

Dashboard ingest URL

Enrollment key name (optional)

☒ Enforce anchor verification (recommended) ?

GENERATE & DOWNLOAD CHARTER

WINDOWS AGENT SIGNING ?

The dashboard holds the agent installer unsigned. Set up your certificate, sign the build, then distribute the signed installer and your CA certificate to your endpoints. The dashboard never keeps a copy of your signing key.

Signing certificate configured **GENERATED**

1

Signing CA created

CN=Self-Sealing Stembolts Corp. Code Signing CA,O=Self-Sealing Stembolts Corp.
CA SHA-256 03f58895d983dc65d770214a...

Your signing key (.pfx) was downloaded when you created it. Vault it; you sign builds with it below.

2

DOWNLOAD AGENT TRUST CERT

Public cert (.crt). Put it on your endpoints (BigFix or Group Policy, into Trusted Publishers) so they trust anything you sign.

[Replace signing CA...](#)

Creates a new CA and signing key and invalidates every build signed with the current one. Use only to roll a compromised key, then remove the old CA from your endpoints.

SIGNABLE AGENT BUILDS

IMPORT UNSIGNED BINARY...

VERSION	PLATFORM	STATE	SIGNED BY	SIGNED AT	ACTIONS
2.8.0.0	linux	ATTESTED	n/a	—	DOWNLOAD
2.8.0.0	windows	SIGNED	CN=Self-Sealing Stembolts Corp. Code Signing,O=Self-Sealing Stembolts Corp.	8/23/2026, 2:33:17 AM	RE-SIGN DOWNLOAD SIGNED DOWNLOAD UNSIGNED

Maintenance

Maintenance configures the server's retention and housekeeping jobs and lets you run them on demand. Data retention sets how long state changes, heartbeats, the audit log, and ingest failures are kept. Endpoint visibility sets the auto-hide window (after so many silent days) and the auto-purge window (after so many hidden days). The BigFix cross-check settings control the nightly enrichment job. A status panel shows row counts, database size, and disk usage.

On-demand action	Effect
Run prune now	Deletes rows past their retention immediately.
Pre-create now	Pre-creates upcoming database partitions.
VACUUM ANALYZE now	Reclaims space and refreshes planner statistics (one to two minutes).
Backfill 168h	Rebuilds the last seven days of trend buckets.

MAINTENANCE ⓘ

DATA RETENTION ⓘ

State-Change Retention
 Days of state-change events and diagnostic reports to retain. Also caps the upper bound of the trend graph.
 days

Heartbeat Retention
 Days of routine "still alive" heartbeats to retain. Kept shorter than state-change retention to bound row counts at scale.
 days

Audit Log Retention
 Days of admin-audit log entries to retain.
 days

Ingest-Failure Retention
 Days of rejected/failed Health Agent submissions to keep for forensic review.
 days

BIGFIX REST CROSS-CHECK ⓘ

Enable BigFix REST Cross-Check
 Pull LastReportTime from the BigFix Root Server REST API to distinguish "Health Agent stopped" from "host went offline".

BigFix Refresh Interval
 How often to refresh the BigFix REST snapshot used for Silent/Disagreement adjudication.
 minutes

Targeted-Query Candidate Cutoff
 Switch the nightly cross-check from a broad environment-wide query to a targeted (filtered) one when candidate count exceeds this number.
 candidates

Large-Environment Threshold
 Total endpoint count above which the nightly cross-check always prefers a targeted query, regardless of candidate count.
 endpoints

ENDPOINT VISIBILITY ⓘ

Auto-Hide After (Days Silent)
 Hide endpoints from every UI surface after this many consecutive days of silence.
 days

Auto-Purge After (Days Hidden)
 Hard-delete hidden endpoints at 03:30 UTC this many days after they were hidden.
 days

HIDDEN TENANTS ⓘ

No tenants currently hidden. (auto-purge is configured at 30 days)

The Hidden Tenants section lists soft-deleted tenants, each with an **Un-hide** or **Purge now** action.

Note: Hidden endpoints are managed on the Endpoints tab (Registered Endpoints, with the visibility filter set to Hidden), not here. This page manages hidden tenants.

HIDDEN TENANTS ?

No tenants currently hidden. (auto-purge is configured at 30 days)

CURRENT STATUS

HEARTBEATS	202 rows	oldest 7/2/2026, 10:41:32 AM
INGEST FAILURES	1,663 rows	oldest 7/29/2026, 11:32:30 PM
AUDIT LOG	206,501 rows	oldest 5/6/2026, 4:27:53 PM
DATABASE SIZE	1227 MB	pg_database_size
DATA DISK	11.6 / 47.9 GB	25.6% used (33.8 GB free)

ON-DEMAND ACTIONS

RETENTION PRUNE

Delete heartbeats older than heartbeat_retention_days and other data older than data_retention_days. Runs automatically every day at 02:00 UTC.

Run prune now

HEARTBEATS PARTITION

Pre-create next month's heartbeats partition so ingest never hits a missing-partition error. Runs automatically on the 25th of each month at 02:30 UTC. Safe to call any time.

Pre-create now

VACUUM ANALYZE

Reclaim space and refresh planner statistics on the high-churn tables (heartbeats, ingest_failures, audit_log, endpoints). Postgres autovacuum normally handles this; use this button after a big retention prune.

VACUUM ANALYZE now

HEALTH-TREND BACKFILL

Recompute the last 7 days of the health-trend pre-aggregation table. Use this if the dashboard's trend chart looks wrong after a long outage. Normal 10-minute aggregator runs handle the rolling window.

Backfill 168h

Maintenance Windows

Maintenance windows suppress expected-outage alerts during planned work. While an endpoint is covered by a window, the dashboard raises no unhealthy, silent, or disagreement alert for it, so a patch cycle or a scheduled reboot does not create a burst of noise. The endpoint's health transitions are still recorded on its timeline; only the alerts are held. Suppression continues for a short grace period (15 minutes by default) after a window closes, so a machine that is still rebooting does not alert the instant it ends.

Two cards at the top of the page show how many windows are active now and how many endpoints are currently in maintenance. To create a window, give it a **Name** and a **Reason** (both required), choose when it runs, add one or more targets, and use **Create window**.

Targets can be any mix of individual endpoints, groups, and tenants. Pick a **Type**, choose the target (for an endpoint, start typing its hostname to search), and use **Add target**; each target appears as a chip you can remove. Targeting a group or a tenant covers every endpoint in it.

MAINTENANCE WINDOWS ?

Planned-downtime windows suppress alerts (unhealthy, silent, and disagreement) for the endpoints they cover. Suppression ends about 15 minutes after a window closes. Maximum window length is 72 hours (change it under Maintenance).

ACTIVE WINDOWS
0

IN MAINTENANCE
0
endpoints

NEW WINDOW

NAME
e.g. October patch cycle

REASON
e.g. Monthly Windows patching

REPEAT
Does not repeat (one-off) ▼

START
08/24/2026, 11:33 PM

DURATION (HOURS)
2

Ends 8/25/2026, 1:33:00 AM (plus 15 min grace).

TARGETS

Add one or more endpoints, groups, or tenants.

TYPE
Tenant ▼

TARGET
PROD Environment ▼

Add target

No targets yet.

Create window

WINDOWS

No active, scheduled, or recent windows.

Use **Repeat** to control how often the window runs:

Repeat	When it runs
Does not repeat	A single window. You set a Start time and a Duration in hours; the page shows the exact end time, including grace, as you fill it in.
Weekly	Every week on a day and time you choose, for the duration you set.
Monthly (day of week)	On an ordinal weekday each month, for example the third Friday, at a chosen time and duration.

For a repeating window you set the day, the time, and the duration; the times are interpreted in your own timezone, which is shown beneath the fields. The dashboard generates each upcoming occurrence automatically, so a recurring window keeps covering its planned downtime without further action.

The WINDOWS list shows each window with its targets, timing, status, and the number of endpoints it currently covers. When a window has several targets, the Targets column shows a count, such as '8 endpoints, 2 groups'; click it to see the full list. Use **End now** to stop a window, or a whole recurring series, immediately, and **Delete** to remove it.

Status	Meaning
ACTIVE	A one-off window running now.
SCHEDULED	A one-off window whose start time is still ahead.
RECURRING	A repeating window; its rule and next occurrence are shown.
ENDED	A past one-off window, kept briefly for reference.
CANCELED	Stopped early with End now, or otherwise canceled.

Note: The maximum length of a single window (72 hours by default) is set on the Maintenance page. The dashboard's 'Active Maint. Windows' card and the 'In Maintenance' status filter on the Endpoints tab both show what is currently suppressed, so a window left open unusually long is easy to notice.

Email/SMTP Settings

Email configures the outbound SMTP relay used for password-reset and notification messages: host, port, encryption (STARTTLS, implicit SSL, or none), username and password, and the From address and name. The stored password is masked; leave it blank to keep the current one. Use **Send test** to send a test message to an address you choose.

EMAIL (SMTP) ⓘ

SMTP HOST	smtp
PORT	587
ENCRYPTION	STARTTLS (port 587) ▼
USERNAME	
PASSWORD	(stored; leave blank to keep current)
FROM ADDRESS	noreply@agentmonitor.com
FROM NAME	Agent Monitor - System Event Notification

Notifications

Notifications holds the global email switches and the alert dispatcher. Global event triggers turn emails on or off for failed logins, account lockouts, password-reset requests, and password changes. System event fallback recipients receive messages that have no more specific destination. Email display timezone sets the timezone shown in alert emails. Alert dispatch is the master switch for alert email (when it is off, alerts still appear on the Priority Alerts tab but no email is sent), with a dispatcher interval that takes effect on the next sweep.

Note: Per-group alert recipients are set on Endpoint Groups, and the silent threshold is set on Health Agent Configuration, not here.

GLOBAL NOTIFICATIONS

GLOBAL EVENT TRIGGERS

Notify on Failed Login
Email the affected user after a failed login attempt on their account.

Enabled

Notify on Account Lockout
Email the affected user when their account is automatically locked.

Enabled

Notify on Password Changed
Email the affected user after their password is changed (reset link, My Account, or a required change).

Enabled

Notify on Password Reset Request
Send the reset-link email when a user requests a self-service password reset. Turning this off disables self-service reset entirely; admins must reset passwords from User Management.

Enabled

SYSTEM EVENT FALLBACK

System Event Fallback Recipients
Leave blank to disable the fallback entirely.

EMAIL DISPLAY

Email Display Timezone
Timezone used to render the Raised-at line of alert emails.

Eastern Time - UTC-04:00 (EDT) - New Yc

ALERT DISPATCH

Alert Email Dispatch
Master switch for the immediate-email alert dispatcher. When off, alerts still land on the Alerts page but no email is sent.

Enabled

Dispatcher Interval
How often the dispatcher sweeps the alerts table for unsent rows. Changes take effect on the next sweep; no restart needed.

5minutes

Backup & Restore

Backup and Restore creates, downloads, restores, and deletes database snapshots, and exports or imports a passphrase-protected server-migration bundle. Use **+ CREATE BACKUP** to take a snapshot; each row offers **Download**, **Restore**, and **Delete**.

Warning: Restoring is destructive: it drops and recreates the database. To confirm, you type the exact backup filename, and the dashboard always takes a pre-restore safety backup first. When a restore finishes it clears your session and you must sign in again.

BACKUP & RESTORE

+ CREATE BACKUP

On-demand `pg_dump` snapshots of `bigfix_health` land under `/opt/bigfix-health-archive/backups/`. Restore is **destructive** (drops and recreates the database); a *pre-restore safety backup* is always taken first so you can roll back if needed.

FILENAME	SIZE	CREATED	ACTIONS
backup_20260821_143909.sql	57.6 MB	8/21/2026, 10:39:11 AM 3d ago	Download Restore... Delete

SERVER MIGRATION

EXPORT FOR MIGRATION

IMPORT MIGRATION

Move this deployment to another server in one passphrase-protected file (`.bch-mig`) carrying the database, BigFix and SMTP credentials, and application configuration. Credentials are decrypted on this server, encrypted with your passphrase for transport, and re-encrypted with the target server's machine key on import. The machine-derived master key itself never leaves either server.

MIGRATION IMPORT INTEGRITY

Every export carries an Ed25519 product signature; the importer verifies it before any decryption attempt so tampered or foreign files are rejected up front. Leave Disabled unless you need to restore a legitimate pre-signing backup whose source you trust.

Allow Unsigned Migration Imports

When Disabled, every `.bch-mig` file must carry a verifying Ed25519 product signature; tampered files are rejected up front before any decryption is attempted. Turn ON only to restore a legitimate legacy unsigned backup whose source you trust; turn back OFF immediately after.

Disabled

About

About shows application and environment information: the product name and version, build date, overall status, and server time, plus database status and size, the configured tenant count, host details (hostname, operating system, Python, CPUs, boot time), disk and memory usage, and the license summary. It is read-only.

ABOUT

APPLICATION

Application	Agent Monitor for BES Clients
Version	2.8.0
Build date	2026-08-05
Overall status	Healthy
Server time (UTC)	2026-08-25T03:35:34.349783+00:00

DATABASE

Status	Connected
Size	1227 MB

BIGFIX TENANTS

Configured	2
------------	---

SERVER

Hostname	agentmonitor
Operating system	Linux 7.0.0-30-generic
Python	3.14.4
CPUs	8
Boot time	8/23/2026, 12:14:13 AM

RESOURCES

Disk usage	11.2 GB / 47.9 GB (24.7%)
Disk free	34.2 GB
Database filesystem	11.2 GB / 47.9 GB (24.7%) on /
Nginx filesystem	11.2 GB / 47.9 GB (24.7%) on /
System memory	15.1 GB (10.1% used)
App memory (RSS)	161.4 MB

LICENSE

Free to use and freely distributable in original unmodified form. Modification for commercial purposes is prohibited.

8.3 Logs

The Logs sub-tab gathers the server's log viewers and audit trails. Each viewer loads on demand, shows the newest entries first, color-codes lines by severity, and exports to CSV (the audit log also exports to PDF). You can filter by level, line count, and a search term.

Viewer	Shows
Application Log	The dashboard service's own log (errors, warnings, and info).
NGINX Log	The web server's error or access log (choose the source).
PostgreSQL Log	The database server's log.
Audit Log	Security-relevant events (who did what, and when), filterable by event type and date.
Ingest Log	Rejected heartbeat and diagnostic posts, with the reason for each.
Performance	A live snapshot of web, database, ingest, system, and network health.

Note: In the access log, entries are classified by HTTP status (a 5xx is an error, a 4xx other than 404 is a warning), not by matching words in the request path, so a normal request whose URL happens to contain a word like 'alert' is not flagged.

Audit Log

Ingest Log

Application Log

NGINX Log

PostgreSQL Log

Performance

SECURITY AUDIT LOG

Login, logout, account changes, setting updates, API key activity, and other security-relevant events.

EVENT

FROM

TO

FILTER

CLEAR

CSV

PDF

Password Changed

mm/dd/yyyy


mm/dd/yyyy

Entries: 10

WHEN	EVENT	BY	TARGET	IP	DETAILS
8/5/2026, 1:51:14 PM	PASSWORD_CHANGED	sansa		192.168.0.169	{"method":"self_service_reset","username":"sansa"}
7/23/2026, 11:45:00 AM	PASSWORD_CHANGED			127.0.0.1	{}
7/23/2026, 11:15:10 AM	PASSWORD_CHANGED			127.0.0.1	{}
7/7/2026, 10:44:26 AM	PASSWORD_CHANGED	sansa		192.168.0.169	{"method":"self_service_reset","username":"sansa"}
7/7/2026, 10:30:32 AM	PASSWORD_CHANGED	mconsuegra		192.168.0.169	{"method":"self_service_reset","username":"mconsuegra"}
7/7/2026, 9:42:52 AM	PASSWORD_CHANGED	mconsuegra		192.168.0.169	{}
7/7/2026, 9:24:17 AM	PASSWORD_CHANGED	mconsuegra		192.168.0.169	{"method":"force_change"}
6/30/2026, 5:01:58 PM	PASSWORD_CHANGED	sansa		192.168.0.169	{"method":"force_change"}
5/13/2026, 2:06:22 AM	PASSWORD_CHANGED	mconsuegra		192.168.0.169	{"method":"force_change"}
5/6/2026, 4:29:21 PM	PASSWORD_CHANGED	admin		192.168.0.169	{}

Performance

Performance is a live snapshot of the dashboard host, covering nginx, PostgreSQL, the ingest pipeline, and system resources. Four gauges (System, Nginx, Database, and Ingest) give an at-a-glance read, with detailed counters below, and the view is cached briefly on the server so multiple administrators can watch it without adding load.



9. Glossary

Terms used throughout this guide.

Health Agent. The small monitoring agent installed on each endpoint. It runs independently of the BES Client, watches the client's health, and reports to the dashboard by exception.

BES Client. The HCL BigFix client whose health the Health Agent monitors. Agent Monitor observes it and never modifies it.

Watchdog. A companion process on each endpoint that reacts immediately when the BES Client service stops, starts, or hangs, so the dashboard learns without waiting for the next scheduled check.

Heartbeat. A periodic report from an agent. Registration and the first report after a reboot also count as heartbeats.

Charter. The signed configuration file delivered beside the installer. It carries the dashboard address, an enrollment key, and the report-encryption anchor, so an endpoint enrolls with no typed parameters.

Anchor. The report-encryption public key and fingerprint assigned to an endpoint at enrollment, so an agent encrypts only to the legitimate dashboard.

Strongbox. Agent Monitor's end-to-end report encryption. Reports are encrypted on the endpoint and can be decrypted only by the dashboard.

Forwarder. A relay node that carries agent reports for endpoints with no direct path to the dashboard. It authenticates with its own key and never reads report contents.

Tenant. A separately scoped BigFix environment within the same dashboard. Access and data can be defined per tenant.

Unhealthy. An endpoint whose BES Client is not reporting normally.

Silent. An endpoint that has not sent a heartbeat within the silent threshold. Silent is distinct from unhealthy, because the endpoint may simply be powered off or off the network.

Disagreement. A mismatch between the two sources of truth: BigFix shows recent activity for an endpoint but the Health Agent reports it silent.

Critical priority. An administrator flag marking an endpoint important enough that its transitions into unhealthy raise alerts and sends notifications.

Build attestation. A vendor-signed statement of a known-good agent build's fingerprint. The dashboard checks each agent against it to detect tampered or unrecognized builds.

Maintenance window. A scheduled period during which alerts for the covered endpoints are suppressed, then resumed automatically when the window ends.

10. Support and Resources

Agent Monitor ships with its full documentation. For first-time setup, see the Installation Guide. To get running quickly, see the Quick Start Guide.

Found a problem or have a suggestion? Open an issue on the project's GitHub repository. Include your dashboard version (shown on the About page), the platform involved (Windows or Linux), and any relevant log excerpt from the Logs tab, so it can be reproduced quickly. The Performance page is a good first stop when the dashboard itself feels slow.